

Multi-Cloud Agentic FDE Ecosystem & Tooling Directory

Generic, AWS, GCP & Azure agentic toolchains mapped to Embedded & Remote FDE PODs

Purpose: A current reference of the agentic-AI tooling landscape for forward-deployed engineering, organized so a practice can pick a stack per cloud and per topology. For each provider it separates **Embedded FDE options** (human-in-the-loop, IDE-native) from **Remote FDE options** (autonomous, managed-runtime).

READ THIS FIRST — A DATING CAVEAT

This directory names specific 2026 products and dates (AgentCore, Kiro, Agent Framework 1.0, and others). The agentic market moves monthly. **Verify every named product and date against the vendor's own docs before quoting this.** The durable layer is the structure (Embedded vs. Remote); the perishable layer is the tool names.

◎ GENERIC & CLOUD-AGNOSTIC ORCHESTRATION

Generic frameworks provide the foundational logic for multi-agent coordination without cloud lock-in — ideal for hybrid deployments where agents span multiple environments or legacy local systems. They are the portable core beneath any cloud-native binding.

Framework	Best for	Embedded fit	Remote fit
LangGraph	Complex, stateful graphs and cyclic workflows.	High — embeds as local state machines in apps.	High — deployable as persistent micro-services.
CrewAI	Role-based collaborative agent squads.	High — ideal for local “coding squads” in IDEs.	Medium — typically needs a managed backend (CrewAI Enterprise).
AutoGen 2.0	Conversation-based multi-agent teams.	Medium — best for interactive chat-based dev.	High — powerful for autonomous, headless loops.
LlamaIndex Workflows	RAG-heavy agents needing massive data reasoning.	High — embedded search & retrieval agents.	High — remote data-syncing agents.
Dify	Visual orchestration and managed platform.	Low — primarily a centralized platform.	High — ready-to-use remote agent infrastructure.

The critical enterprise choice

The decision that matters most for adoption is **managed platform vs. developer framework**: platforms (Dify, Nexus-class) trade control for speed-to-deploy; frameworks (LangGraph, CrewAI) trade setup effort for portability and control. Regulated/air-gapped buyers usually need the framework path; fast-moving product teams often prefer the platform path.

AWS AGENTIC ECOSYSTEM (June 2026)

Per the source: following the June 2026 AWS NY Summit, AWS shifted from “Copilots” to “AgentCore” infrastructure — positioning agents as a combination of models plus managed harnesses. *[Verify product names/dates before use.]*

Embedded FDE options

- **Kiro** — a developer toolchain with a mobile app and IDE integration for validating infrastructure and testing.
- **AWS DevOps Agent** — automated release-readiness and remediation, working directly inside CI/CD pipelines.
- **Strands Agents SDK (v1.0)** — open-source toolkit for production agents with integrated chaos testing and context management.

Remote FDE options

- **AgentCore Harness** — a managed execution environment that decouples agent logic from models; runs agents on isolated micro-VMs, handling state persistence and error recovery autonomously.
- **AWS Context Service** — unified knowledge-graph construction so remote agents maintain consistent codebase topology.
- **Amazon Q Enterprise** — an enterprise-wide AI assistant deployable as an autonomous service for codebase modernization.

Performance & infrastructure

- **Compute:** EC2 G7 instances with NVIDIA Blackwell GPUs for high-speed agent inference.
- **Retrieval:** Web Search on Amazon Bedrock AgentCore.

GOOGLE CLOUD (GCP) AGENTIC ECOSYSTEM

Per the source: GCP consolidated its offerings into the Gemini Enterprise Agent Platform (formerly Vertex AI), focusing on long-running autonomous agents and multimodal reasoning. *[Verify product names/dates before use.]*

Embedded FDE options

- **Google ADK** — the native framework for building agents on Gemini; stable v1.0 across four languages in 2026.
- **Workspace Studio** — a no-code agent builder embedded in Google Workspace for process-specific business agents.
- **Agentic Vision** — a Gemini 3 Flash capability letting embedded agents process visual development assets in real time.

Remote FDE options

- **Gemini Enterprise Agent Platform** — supports long-running agents that independently solve complex problems over several days.
- **Agent Engine (Agent Runtime)** — a managed service providing secure code execution, sessions, and “Memory Banks” for remote agents (billing live Feb 11, 2026 per source).
- **A2A Protocol v1.0** — an agent-to-agent communication standard, reported in production at 150+ organizations for multi-cloud interoperability.

Performance & infrastructure

- **Compute:** TPU 8t (training) and TPU 8i (inference), engineered for the agentic era.

MICROSOFT AZURE AGENTIC ECOSYSTEM

Per the source: Microsoft launched Agent Framework 1.0 (April 2026), the production-ready convergence of Semantic Kernel and AutoGen into a single unified SDK. *[Verify product names/dates before use.]*

Embedded FDE options

- **Agent Framework 1.0** — native .NET and Python support with stable APIs for embedding agentic logic into enterprise apps.
- **Microsoft Foundry VS Code Extension** — GA as of June 3, 2026; deploy hosted agents and browse 1,900+ models in the editor.
- **Azure Cosmos DB Agent Memory Toolkit** — standardizes persistent agent memory using local emulators for dev/test without cloud dependency.

Remote FDE options

- **Azure AI Agent Service** — a managed service for multi-agent systems, integrating with Foundry for evaluations and memory stores.
- **Foundry Agent Service (BYOM)** — “Bring Your Own Model” — remote agents can use fine-tuned or third-party models behind Azure API Management.
- **Azure Container Apps Dynamic Sessions** — specialized compute for fast, secure isolation of remote agent code execution.

The adoption pattern across every cloud is the same: deploy **Embedded** for trust-building and developer speed, and **Remote** for high-scale autonomous service delivery. This crosswalk maps the two topologies across personas, interfaces, tools, and gates.

Component	Embedded PODs (Human-in-the-Loop)	Remote PODs (Autonomous Services)
User persona	Individual developers, team leads.	DevOps engineers, SREs, product owners.
Primary interface	IDE (VS Code, Cursor), Slack, GitHub PRs.	API gateway, event bus (Kafka/SQS).
Core tools	Kiro, Foundry VS Code, Strands SDK.	Bedrock AgentCore, GCP Agent Runtime, AutoGen.
Orchestration	LangGraph, Semantic Kernel.	Azure AI Agent Service, Gemini Platform.
Acceptance criteria	Human code review / approval gate.	Automated evaluations / canary deployments.

Mapping the directory to the 80:20 model

This crosswalk is the tooling expression of the 80:20 split. **Embedded** = ~70:30 (the human holds the approval gate in the IDE); **Remote** = ~90:10 (automated evaluations and canary deploys hold the gate, under human-set policy). Same accountability, different interaction surface — exactly the topology split defined in the FDE 80:20 lifecycle.

Industry trends (June 2026, per source)

- **Agentic SDLC:** over 70% of teams report high GenAI usage across the full development lifecycle.
- **Talent shift:** entry-level coding roles are shifting toward senior-agent collaboration — agents perform tasks formerly assigned to junior engineers.
- **Managed vs. framework:** the distinction between managed platforms and developer frameworks is the most critical enterprise adoption decision.

HOW TO USE THIS DIRECTORY

This directory is a living document, updated quarterly. It is not a static list of tools, but a dynamic map of the ecosystem. The information is organized by cloud provider, then by tool type, and finally by specific tool. The goal is to help you find the right tool for your needs, and to understand the relationships between different tools and providers. The directory is designed to be easy to navigate, with clear links to the tools and their documentation. It is a resource for developers, architects, and anyone interested in the multi-cloud agentic FDE ecosystem.